



THE UNIVERSITY OF AUCKLAND

COMPSCI 780

Benchmarking software obfuscators An experimental approach

Rafik Nasraoui

supervised by
Dr. Clark THOMBORSON

May 29, 2016

Abstract

A high level of security, to the extent that all hacks are rendered impossible, is of crucial importance in most software. Obfuscation is a potential method of protecting sensitive information, whereby the computer program is rendered unintelligible to a human being, without altering the core functionality of the program. As in any secure system, obfuscators increase the work an attacker needs to perform, however a potential tradeoff in the use of obfuscation is that it may impact the execution speed of a program, especially if the program size exhausts the CPU cache capacity, and additional calls to main memory are required. To test the relationship between obfuscation and execution speed, we modelled a benchmark program using a fixed-width boolean circuit, where the width of the circuit is increased with each execution. We found that execution speed decreases as circuit width increases, with a marked increase in the execution time when the circuit width becomes wider than the CPU's highest cache level. However, the latency increase due to the shift to main memory is markedly less than was anticipated, the causes of which need to be further investigated.

1 Introduction

Security and data privacy are critical aspects of most software programmes. The need for such security can derive from the need for compliance with government regulations, and individual and company privacy requirements. The protection of sensitive information and industry data is also vital for software to stay competitive. The ability of hackers to extract sensitive data is continuously advancing, and so software data protection must also continuously evolve, in order to keep ahead of such efforts.

One of the mechanisms that could be implemented in the future to protect data is obfuscation. Software obfuscation aims at making computer programs unintelligible for a human being, without affecting the functionality of the software. Although software must be delivered into the hands of the public in order for the public to make use of the software, obfuscation hides the information contained within the software, including the code itself, by rendering it illegible to a user or hacker.

Obfuscation has potential useful applications in preventing reverse engineering and thereby protecting Intellectual Property, and in enabling secure software distribution. Obfuscation can also be a powerful means of encrypting data.

Until recently, there was no formal definition of obfuscation, and all techniques, including manually obscuring data, could be given the label obfuscation. These techniques do not have quantifiable security and anyone with appropriate, yet easily available, tools such as a debugger and a compiler can reverse engineer code that has been obfuscated with the current best efforts[9].

Recent developments in cryptography theory have however produced a formal definition of obfuscation, which will most assuredly lead to an increased interest in the field, in the hope that obfuscation can be practically applied to protect software in the near future. The theory also suggests that the adversary work factor, i.e. the work required to be performed by an adversary to break the obfuscation, scales exponentially with the program runtime.[9]

However, the drawback of using obfuscation is that it may be ultimately limited by the hardware on which the software is run. Obfuscators need to produce obfuscated programs using available computing resources in a commercially relevant amount of time. These obfuscations then need to be evaluated by end users in relatively similar, if not more challenging conditions, in terms of hardware performance and functionality. Normal encryption methods in general have small variable size and a low number of instructions, and therefore require little run time to execute. Obfuscators however are technically challenging to implement, and the number of variables and instructions arising from an obfuscation may surpass the limits of the hardware within which the software is run.

This paper aims to test the hypothesis that obfuscation is limited by the hardware it runs on. We argue that the run time of a boolean gate that belongs to an obfuscated circuit can be represented by the following model:

Should this hypothesis prove true, it would place serious constraints on the practicality of obfuscation in most situations.

2 Background

This section outlines the definitions of boolean circuits and indistinguishability obfuscation that are essential to this paper.

2.1 Boolean Circuits

Real world programs are compiled into a sequence of machine instructions that can be interpreted by a processing unit (CPU) manipulating a finite number of *bits* at a time. The CPU is a chip that is made of logical gates that are made of a few transistors. In this context, boolean circuits are a powerful model that can express any algorithm we can run on a CPU[4].

Complexity Theory defines a *Boolean Circuit* C with n inputs and m outputs as a finite, labelled, directed acyclic graph. It contains n nodes with no incoming edges or ancestors; called the input nodes and m nodes with no outgoing edges or successors, called the output nodes. All other nodes are called gates and represent the logic operation AND, OR and NOT[2]. Simply put, a *Boolean Circuit* is a diagram showing a combination of logic operators to drive an output from an input. By associating a boolean function with each gate, a boolean circuit computes an arbitrary function

$f \in F : \{0,1\}^n \rightarrow \{0,1\}^m$. Two common properties are associated with boolean circuits:

1. The *size* of a circuit is the number of non-input vertices that it contains.
2. The *depth* of a circuit is the length of the longest directed path from an input vertex to an output vertex.

Because of the acyclic and oriented nature of the circuit, each node can be represented as receiving its inputs from a lower order depth only. To evaluate a circuit, we process the nodes in an increasing order of depth by applying the associated boolean operation on the input vertices. The result is passed to the higher order nodes.




Figure 1: Simple high level representation of a boolean circuit

Size and Depth give an indication of the computation necessary to evaluate a circuit. Size is a rough measure of the computation time a single processor would need for a sequential circuit evaluation. Depth is an important cost metric for time efficiency in a parallel execution environment where the size is no longer the main bottleneck[13].

Another established circuit property that is relevant to our study is the *width* of a circuit[12]. If the circuit is arranged in levels such that the input of each level are connected from lower order levels exclusively, we define the width w as the maximum number of gates contained in single level.

The formal study of circuit complexity is beyond the scope of this paper. We should however quickly mention the notion of *Circuit Family*. Our definition limits the size of input to a fixed number of bits. The notion of boolean circuit C_n computing f_n is generalised to a family of circuits C which is a collection of boolean circuits $\{C_1, C_2, C_3, \dots\}$ such that there is one circuit C_i that computes F on i inputs.

2.2 Indistinguishability Obfuscation

Obfuscation must meet two requirements, as defined by Barak et al.[14]:

1. The functionality of the obfuscated program must remain identical to the functionality of the original, un-obfuscated program.
2. Anything (and everything) that can be efficiently computed from O_C can be efficiently computed given oracle access to C . i.e. O_C does not disclose any information about the input other than what C does.

The first requirement is trivial and means that for any input program C , an Obfuscator O produces a new program O_C where the functionality of O_C is identical to the functionality of C . This condition means that given an input space I we have $\forall x \in I, C(x) = O_C(x)$.

The second requirement is less trivial and also known as the Virtual Black Box (*VBB*) requirement. It implies that there is no ‘leakage’ of information, other than the input and the output. This means that it is impossible to tell whether an output was produced from the computation of C , or from the computation of O_C . An adversary can learn a predicate of $O_C(x)$ with the same probability of learning it from $C(x)$.

Barak et al. concluded that the VBB requirement is impossible to attain in certain situations[14], however, it is possible to come close to meeting the second requirement, which in practice is enough to maintain the usefulness of obfuscation as an effective security measure. This is done by modifying the requirement of VBB, and replacing it with two weaker definitions that can still result in rendering a program unintelligible in a meaningful and practical way. These definitions are namely indistinguishability *iO* and differing-input obfuscation *diO*.

iO requires that for any circuit pair C_1 and C_2 from the same circuit family C that compute the same functionality f , the obfuscated versions of C_1 and C_2 ; $iO(C_1)$ and $iO(C_2)$ are perfectly indistinguishable. The probability of an observer being able to tell if an output was generated from $iO(C_1)$ or from $iO(C_2)$ is negligible.

The stronger notion of differing-inputs obfuscation *diO* implies that given a circuit pair (C_1, C_2) , if it is hard to find x such that $C_1(x) \neq C_2(x)$ then it is hard to distinguish between $diO(C_1)$ and $diO(C_2)$ [8]. The existence of an adversary that can distinguish between obfuscations of circuits C_1 and C_2 implies the existence of an adversary that can find an input on which the two circuits differ[14].

Some work has been done within the scientific community to develop real-life obfuscations that meet the requirements given by Barak et al. Garg et al.[7] published an obfuscation construction valid for all polynomial-size, log-depth circuits constructed under novel algebraic hardness assumptions. Based on this work, Banescu et al.[3] published a proof-of-concept implementation of this *iO* construction. However, to date there is no efficient practical implementation of Barak et al.[14] indistinguishability obfuscation requirement. The study of the security of functional obfuscators is beyond the scope of this paper.

But this lack of ready-made obfuscators does not limit the ability of this paper to prove or disprove the hypothesis that obfuscation is ultimately limited by the hardware restraints of the operating system, as a simulation of an obfuscation-like process is sufficient for this purpose.

3 Hypothesis

We suggest that, for on a fixed hardware platform, the total running time of a boolean circuit follows this analytical model:

$$T = \max\{t_g + t_e, t_m(w)\} \quad (1)$$

where T is the total elapsed (wallclock) time in processing a single gate, t_g is the CPU time for generating a gate, t_e is the CPU time for evaluating a gate, and $t_m(w)$ are the memory-stalls from the gate-generation and gate-evaluation computations.

Any hardware platform evaluating an obfuscated circuit C needs to represent its gates in a readable and comprehensible internal format that is stored in memory, then apply the logical function of every gate on the gate's relevant inputs and finally store those results in memory to be used by other gates. The evaluation of each gate has a time cost t_g related to the creation of its representation on the evaluation's platform, and a second time cost t_e related to the pure evaluation of the logical function of the gate.

Our prediction is that t_g and t_e are constant or of a similar order for any gates in any given random circuit of the same width w .

We introduce $t_m(w)$ as the variable time component in a gate's evaluation. It is related to the organisation of the circuit and how it is represented on a platform's memory. It can also vary from one platform to another. But for a fixed platform, we predict that it is a function of the width w . If the width w is wider than the last level of cache of CPU of the hardware platform evaluating the obfuscated circuit under test, then we are most likely to experience memory stalls that severely affect performance. In that scenario, the evaluation time is dominated by the time the CPU spends fetching the inputs of a gate, as opposed to the time it requires to evaluate the gate's logical function, once the inputs are available. This would confirm that the overhead introduced by obfuscation makes obfuscated program unadapted to current hardware and therefore inefficient for end users.

If the proposed analytical model is validated by our results, we expect it to be also valid for other platforms of a similar range.

4 Methodology

We now describe the algorithm used to simulate an obfuscated boolean circuit and the tools used to measure its execution time on a hardware platform.

4.1 BPW Generation

Thomborson proposed a space-efficient and computationally-appropriate file format called *BPW* for the evaluation of very large Boolean functions with

bounded program widths[12]. In the *BPW*, Gates of a circuit are represented by descriptors holding information about their type (AND, OR, ...) and are sequentially encoded inside the file's body. In our experiment we generate random boolean circuits (which evaluate random boolean functions) with varying width w for a fixed circuit size N . The relationship between a circuit width w and depth D is $N = w \times D$, so for $w = N$ the circuit has depth 1.

The generated circuits are represented using a simplified BPW format by storing the gate type and the input indices only and by making following assumptions about the properties of the circuit under test.

1. All circuits have size $N = 10^8$
2. w takes the values $\{10^3, 10^4, 10^5, 10^6, 10^7, 10^8\}$
3. A circuit input array has size w .
4. Any given level L has exactly w gates.
5. Every gate has only one output.
6. A circuit output array has size w .
7. Any given gate from a level L_i has between one and three inputs that come from level L_{i-1} exclusively, such that: $G_i = f(X)$, where f is a boolean operator and X is an 8-byte array of size 1, 2 or 3.
8. The output of a gate is written inside the output array at the same index of the gate. $G_i(X) = \text{Output array}[i]$.

The *rand()* method from the *C* standard library was used to attribute a random type to a gate, as well as assign its input indices.

The entire circuit is generated prior to its evaluation. The generation step executed initially is therefore not monitored for performance. The resulting circuit is stacked gate by gate in a simplified BPW file.

Following the initial BPW format developed by Thomborson, the first 4 bytes represent the header of the a BPW file[12]. The following bytes then represent the total gates of the circuit. All gates are sequentially represented from index 0 to index $10^8 - 1$. A given gate's encoding varies from 9 to 25 bytes. The first byte holds the type of the gate. The following 8 to 24 bytes present the input indices encoded over 64-bits each. The gate's representation could be optimised to the minimum bytes needed to encode a gate. Compression could also be used to save disk space. We chose to keep the existing model for simplicity. As mentioned earlier, the size of the circuits under test was fixed to $N = 10^8$. We modelled 14 types of Gates; 1×1 -input, 6×2 -input and 7×3 -input gates. On average, a

gate has 2.43 inputs encoded with 8 bytes per input plus one byte for the gate type. The average size of a BPW file storing a circuit is therefore $avg(sizeof(Gate)) * 10^8 \approx 2GB$.

4.2 Evaluation

A circuit evaluation is sequential for both gates and levels. During execution, gates of a level L_i are processed in ascending order of index in the level. Inputs indices for a gate are randomly distributed over $[0, w-1]$. Every gate retrieves its relevant inputs as described by the gate descriptor. Once a gate is processed, its output is stored in the output array at the same index as the gate. Once all the gates of level L_i are processed, the execution cursor moves to level L_{i+1} . The hypothesis is that the random distribution of gate input indices will significantly affect the overall performance of the circuit for large w , as the entire input array cannot fit in the cache. Figure 3 describes how a level is evaluated.

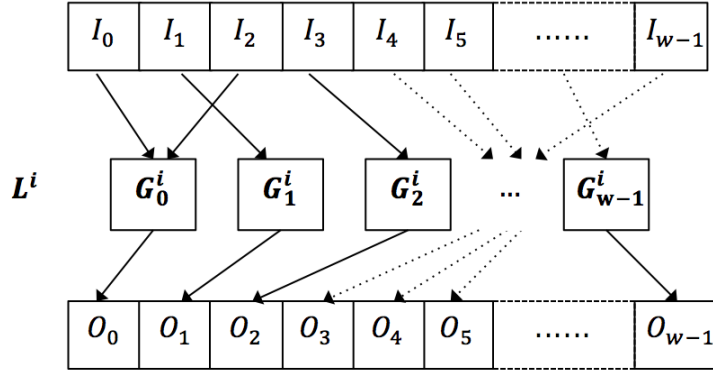


Figure 2: Simple description of Level L_i evaluation

4.3 Hardware

Obfuscators might occasionally run on super computers, however it is likely that viable commercial use of obfuscators would require them to run on mass market platforms such as smart phones, battery powered laptops and desktop computers[12]. The target platform for the circuits evaluation is a mid-level desktop computer. The CPU is a 64-bits-instruction-set with a base frequency of 3.40 GHz (Max of 4.1GHz). It has 4 cores that can run a thread each and has the following cache capacities (in bytes):

- Level 1 data cache: 32K
- Level 1 instruction cache: 32K
- Level 2 cache: 256K

- Level 3 cache: 6144K

We predict that having a circuit with a width higher than the capacity of the last-level cache is a perfect candidate to observe an impact on performance.

The operating system used is Linux Fedora 22 on kernel 4.4.5.

4.4 Metrics

The *perf* linux command was used to measure the impact of evaluating a circuit on the CPU. *perf* is a tool written in C that uses Performance Counters to profile an application ¹. Performance counters are CPU hardware registers that count hardware events such as instructions executed, cache-misses suffered, or branches mispredicted[11]. The *perf* tool has the advantage of being program oriented, instead of system oriented [1]. In particular, the *perf stat* subcommand allows to retrieval event counts that are of interest to our experiment. These metrics are mainly *instructions*, *cycles* and *LLC-loads*.

- **instructions** are unitary operations executed by the CPU at the lowest level such ADD, SUB, LOAD or STORE. They can represent arithmetic operations, data handling operations or control flow operations. Modern CPUs include more complex instructions for complicated integer or float point operations.
- **cycles** are the steps performed by the CPU when it receives an instruction. Typically early processors would pipeline fetching the instruction from the program counter, decoding it, reading its associated variables from the registers, executing its logic order and finally writing the result back to registers or storing it the system's main memory[10]. After an initial set up time, pipelining allows to execute one instruction per cycle, assuming no stalls. Modern CPUs achieve more than one instruction per cycle using techniques like *Instruction-level Parallelism* (loading more than one instruction in the pipeline) and *out-of-order execution* (allowing the simultaneous execution of multiple instructions that don't depend on each other for input or output). The Intel Core i5 can handle up to 4 instructions per clock cycle[6], but this figure is usually lower because of stalls, i.e. when the pipeline needs to for the result of a previous instruction before proceeding.
- **LLC-loads** are instructions to load results from the Last-Level of Cache (or L3) of the CPU. L3 cache is the largest and highest-latency level of cache for the Intel Core i5 CPU. It is shared amongst all its cores.

¹<https://github.com/torvalds/linux/tree/master/tools/perf>

5 Simulation Results

In our hypothesis formulated by equation 1, we argued that the total runtime of a gate is either dominated by the deterministic time of generation and execution of the gate when w is small, or by the variable time the CPU spends preparing the inputs for a gate’s evaluation when w is large.

For every pre-determined value of w , *perf* runs the experiment 3 times and collects the arithmetic average of the event counters. The data collected is presented in tables 1 and 2.

Table 1: Average results of 3 simulation runs for $w \in \{10^3, 10^4, 10^5\}$.

w	10^3	10^4	10^5
task-clock (ms)	13116.554712	13322.905953	13313.335011
instructions	99200423097	97602445259	97553743526
cycles	55393862507	56255744084	56207517444
instruction/cycle	1.74	1.76	1.74
branches	17179168771	16896966649	16887492840
branch-misses	213275235	267635214	267596585
LLC-loads	75319029	75177952	78311853
std deviation	1.49%	0.94%	0.27%

Table 2: Average results of 3 simulation runs for $w \in \{10^6, 10^7, 10^8\}$

w	10^6	10^7	10^8
task-clock (ms)	13633.97763	16041.15891	17507.07431
instructions	97578653791	98201779911	103646369628
cycles	57584456723	67736311966	73941409511
instruction/cycle	1.69	1.44	1.4
branches	16897296695	17058902603	18501020880
branch-misses	266970610	266910904	268274635
LLC-loads	328139283	380853849	505923626
std deviation	0.43%	0.50%	0.19%

From the results collected in tables 1 and 2 we assume that $t_m(w)$ for $w = 10^3$ is negligible. In our experiment $E(w = 10^6)$ the total run time is assumed to reflect load and execution operations only with no significant width penalty. For $w = 10^3, N = 10^8$, we have $N \times T = 13116\text{ms} \implies T = t_g + t_e \approx 130\text{ns}$.

For $w > 10^7$ the computation becomes memory-bound, allowing us to estimate $T = 162\text{ns}$ and $t_m(10^8) = 46\text{ns}$. Our initial results seem to suggest that

$$T = \begin{cases} t_g + t_e & \text{if } \log_{10}(w) < 7 \\ t_g + t_e + t_m(w) & \text{if } \log_{10}(w) > 7 \end{cases}$$

which disproves the hypothesis stated in equation 1 as $t_m(w)$ is not dominant compared with $t_g + t_e$.

Since the CPU has a frequency of 3.40 GHz the load and execution of a gate require about 450 CPU cycles per gate, which is up to 2000 instruction per cycle (at 1.7 instructions per cycle) on a quad-core CPU assuming no stalls. For all tested $w < 10^7$ we observed a stable instruction per cycle rate at 1.7 ins/cycle with 0.2 instruction per cycle stalled. By dividing the the total instructions observed by the total number of gates, we establish the instructions per gate at around 1000 instructions. Which suggests that the CPU was running at about 4Ghz, a higher frequency that the base 3.4 Ghz frequency.

Table 3: Average simulation results per Gate

width	time/Gate (ns)	LLC-loads/Gate	instructions/Gate
10^8	175.0707431	5.05923626	1036.463696
10^7	160.4115891	3.80853849	982.0177991
10^6	136.3397763	3.28139283	975.7865379
10^5	133.1333501	0.78311853	975.5374353
10^4	133.2290595	0.75177952	976.0244526
10^3	131.1655471	0.75319029	992.004231

While w increases, the instructions per gate average varied slightly between 976 and 1036 instructions i.e. a standard deviation of 24 instructions, or about 2%. However, the task-clock time average spent per clock has increased from $130ns$ for $w = 10^3$ to $175ns$ for $w = 10^8$. The standard deviation observed is about 12%, or about $18ns\%$.




Figure 3: Evolution of task-clock time per gate in function of width w . A base-10 log scale is used for the X axis. The trendline is $t(w) = 0.67 \log_{10}(w) + 126.52$, $R^2 = 0.86748$ based on the data collected for $w \leq 10^6$. The 5% confidence interval boundaries are plotted in red.

Since we used six types of 2-input gate and seven types of 3-input gate (out of a total of fourteen gate types), a gate evaluation has one fetch to read the gate type and about 2.5 memory-fetches on average to read its inputs. The gate evaluation process also writes the result in the output buffer which causes additional memory reads and writes.

Once the gate is unpacked in a top-level *C struct* of 16 bytes (1 byte for the type, 8 bytes for the input pointer, before padding to the largest power of 2), with 2.5 input indices of 8 bytes each, a level needs $(16 + 2.5 \times 8) \times w$ of memory to be fully represented, e.g. levels for $w = 10^7$ need 3.6 GB. This amount of memory cannot obviously reside in cache, and on 4 GB memory platforms there would certainly be page-faults requiring hard disk activity. We estimate that the level array can fit inside our system’s 16GB RAM without disk faults. Our initial algorithm loads an entire level before feeding into the execution loop. We could improve this by evaluating a gate as soon as it is loaded.

We designed the gate inputs to be spread randomly in an array of width w . For any large enough w , the CPU last level cache (LLC) will miss with a high probability when the input bit-array cannot be fully contained. On our test platform, LLC is 6MB wide and the cache line is 64B[6].

For $w = 10^7$, the input array occupies 8×10^7 bits, the miss rate should be $1 - (6 \times 10^6 B / 10^7 B) = 40\%$. From the observed $T = 162ns$ we deduce the memory latency, which is the time required by the CPU to fetch a gate input from the cache, to be $162ns$ per LLC miss.² Each unpacked level occupies 36×10^7 bytes in memory. For the observed run time this is equivalent to 225 MB/s in memory bandwidth. The reported LLC-loads represent about 10% for the bandwidth at 23.5 MB/s. However, the cache line being 64B wide, the LLC-loads represent 1.3 GB/s bandwidth. With the estimated latency, there should be a gate input miss every 162ns. This gives us a bandwidth of 453 MB. We suspect that the 900 MB difference noted is possibly due to the read and write traffic caused by the writing of the results back into the output array.

The Ivy Bridge CPU family, to which our test CPU belongs, also seems to have a problem with the pre-fetch instructions wasting time on pre-fetching data that are already in the cache[6] which might explain the high bandwidth observed.

Based on the data on this paper, we could not observe significant hit on the overall performance when evaluating wide circuits. The average LLC-loads per gate do behave as predicted, as shown in 3, where a large increase is seen around the $w = 10^6$ mark. The LLC-load rate was almost constant for $w < 10^5$ at 0.75 load per gate but then jumped 5 loads per gate for $w = 10^8$. However, the increase in the LLC-load rate did not correspond with a similar increase in execution time.

² $latency = \frac{T}{2.5 \times miss\ rate}$ where T is the total run time per gate

This might be explained by the CPU’s successful speculative optimisation of branch mispredictions. We also suspect that the generation of random 64-bit input indices using the C `rand()` method (which returns a random 32-bit integer) might produce numbers with the most significant bits 75% biased towards 1³. This will affect the memory-fetch statistics for gates with large indices.

6 Discussion and Future Work

Obfuscation is a potentially powerful tool to encrypt and protect data in the future. However, one of the limitations of obfuscation may prove to be that it lowers the execution speed of a programme, at least in some circumstances.

In this experiment we simulated an obfuscated circuit by generating a fixed-width random boolean circuit and evaluated the performance of that circuit. We have shown how *perf* hardware counters can be used to measure cache hits inside a 6MB L3-cache ix86 microprocessor clocked at 4GHz in order to evaluate a simulated obfuscated circuit.

We have proposed an analytical model to explain the execution times observed. Data collected for width $w > 10^6$ falls outside the confidence interval of the regression curve of smaller widths. Further statistical analysis will be done before this work will be published in a technical journal. In this study our assumption was that when w is large, it would be possible to separate $t_g + t_e$ and t_m , as $t_g + t_e$ would become negligible as $t_m(w)$ would increase significantly and dominate the results. However the results obtained failed to match this prediction

In future work, we plan to optimise the code so that a gate evaluation would be performed in tens of CPU cycles only. Although our CPU has a maximum specification of 8 instructions per cycle^{??}, we think that attaining that ceiling will be quite hard since most of those instructions are low-level instructions needed internally by the CPU.

We hope however that by reducing t_e by a factor 10 we can better evaluate $t_m(w)$ and t_g . Our current experiment was not optimised for t_g or $t_m(w)$. By generating the entire circuit prior to evaluation, the computation was either memory bound, file-IO bound, or stalls and cache-conflict bound when the output of the generator process was piped to the input of the evaluator process. We can also improve the size of program in memory by encoding input and output values over 1 bit rather than the current 1 byte encoding. We believe that there is a trade off between memory and CPU when evaluating t_g . When not reading from a file descriptors, gates would be created using random number generators that would require extra CPU cycles.

³<http://stackoverflow.com/questions/4945698/is-the-value-of-rand-max-always-2n-1>

We don't think that designing a multithreaded experiment will improve the observed running speed if the memory bandwidth of the CPU is already saturated and that $t(w)$ can only be lowered by reducing the memory loads to a minimum.

We are also focusing on improving our simulation by replacing the random generation mechanism using the *C* standard library *rand()* method with the stride approach[5] to remove any bias in generating large random integers and break any predictability patterns.

Future work would explore the use of the Intel Performance Counter Monitor library⁴ as a high level interface for measuring real time performance data directly inside the circuit's code to separate . In particular, we plan to measure the CPU cache and instruction metrics at initialisation time, then at generation time and finally at execution time to isolate each stage's impact on performance.

A more extensive study, using a large range of hardware devices and over a large range of circuit simulations, could also help validate or disprove the proposed model.

References

- [1] R. Arnold. A. Zanella. *Evaluate performance for Linux on POWER. Analyze performance using Linux tools*. 2012. URL: <http://www.ibm.com/developerworks/linux/library/l-evaluatelinxonpower/>.
- [2] Sanjeev Arora and Boaz Barak. *Computational Complexity: A Modern Approach*. Cambridge University Press, 2009, pp. 101–103. URL: <http://theory.cs.princeton.edu/complexity/circuitschap.pdf>.
- [3] Sebastian Banescu et al. “Idea: Benchmarking Indistinguishability Obfuscation – A Candidate Implementation”. English. In: *Engineering Secure Software and Systems*. Ed. by Frank Piessens, Juan Caballero, and Nataliia Bielova. Vol. 8978. Lecture Notes in Computer Science. Springer International Publishing, 2015, pp. 149–156. ISBN: 978-3-319-15617-0. DOI: 10.1007/978-3-319-15618-7_12. URL: http://dx.doi.org/10.1007/978-3-319-15618-7_12.
- [4] Stephan Mertens Cristopher Moore. *The Nature of Computation*. OUP Oxford, 2011, p. 131.
- [5] Clark Thomborson Enyou Li. *Data Cache Parameter Measurements*.
- [6] Agner Fog. *The microarchitecture of Intel, AMD and VIA CPUs An optimization guide for assembly programmers and compiler makers*. 2016. URL: <http://www.agner.org/optimize/microarchitecture.pdf>.

⁴<https://software.intel.com/en-us/articles/intel-performance-counter-monitor>

- [7] Sanjam Garg et al. *Candidate Indistinguishability Obfuscation and Functional Encryption for all circuits*. Cryptology ePrint Archive, Report 2013/451. 2013.
- [8] Sanjam Garg et al. *On the Implausibility of Differing-Inputs Obfuscation and Extractable Witness Encryption with Auxiliary Input*. Cryptology ePrint Archive, Report 2013/860. <http://eprint.iacr.org/>. 2013.
- [9] Michael Hsieh. *SAFEWARE, Program Information*. Defense Advanced Research Projects Agency. <http://www.darpa.mil/program/safeware>.
- [10] John Morris. *Computer Architecture*. SOFTENG363, University of Auckland. 2008. URL: <https://www.cs.auckland.ac.nz/~jmor159/363/html/pipelines.html>.
- [11] Perf. *Linux Kernal Organisation - Perf Wiki*. 2015. URL: <http://perf.wiki.kernal.org>.
- [12] Clark Thomborson. “Benchmarking Obfuscators of Functionality”. In: *CoRR* abs/1501.02885 (2015). URL: <http://arxiv.org/abs/1501.02885>.
- [13] Ryan Williams. *An Overview of Circuit Complexity*. URL: <http://web.stanford.edu/~rrwill/week1.pdf>.
- [14] B. Barak O. Goldreich R. Impagliazzo S. Rudich A. Sahai S. Vadhan K. Yang. *On the (Im)possibility of Obfuscating Programs*. 2001.